

Operational Transparency: Showing When Work Gets Done

Robert L. Bray

Kellogg School of Management, Northwestern University

April 17, 2019

Abstract

Problem: Do the benefits of operational transparency depend on when the work is done?

Methodology: I study how customers respond to operational transparency with parcel delivery data from Cainiao Network, the logistics arm of Alibaba. The sample describes 4.68 million deliveries. Each delivery has between four and ten track-package activities, which customers can check in real time, and a delivery service score, which customers leave after receiving the package. I regress the delivery scores on the track-package-activity times.

Results: I show that customers punish early idleness less than late idleness, leaving higher delivery service scores when track-package activities cluster toward the end of the shipping horizon. For example, if a shipment lasts 100 hours, then delaying the time of the average action from hour 20 to hour 80 increases the expected delivery score by approximately the same amount as expediting the arrival time from hour 100 to hour 73.

Relevance: I bridge OM's literature on operational transparency with psychology's literature on the peak-end effect.

Managerial implications: Memory limitations make customers especially sensitive to how service operations end.

Keywords: Operational transparency; package delivery; peak-end rule; empirical operations management.

1 Introduction

Customers rate service more highly when effort is visible. For example, Buell and Norton (2011, p. 1564) argued that “engaging in operational transparency, by making the work that a website is purportedly doing more salient, leads consumers to value that service more highly.” And Buell et al. (2017, p. 1673) reported that

The introduction of [operational] transparency contributed to a 22.2% increase in customer-reported quality and reduced throughput times by 19.2%. Laboratory studies revealed that customers who observed process transparency perceived greater employee effort and thus were more appreciative of the employees and valued the service more.

These authors, however, implicitly assume that customers will perceive worker effort, but not worker loafing. Yet true operational transparency will make both industry and idleness visible. For example, consider the following track-package records that Cainiao Network shared with its customers:

Shipment	Action	Facility	Timestamp	Score
3144672	Order	-	2017-02-05 15:05	-
3144672	Consign	-	2017-02-05 17:37	-
3144672	Receive	105638	2017-02-05 18:40	-
3144672	Depart	105638	2017-02-05 21:52	-
3144672	Arrive	65132	2017-02-06 04:15	-
3144672	Depart	65132	2017-02-06 05:20	-
3144672	Arrive	29048	2017-02-06 08:22	-
3144672	Scan	29048	2017-02-06 08:44	-
3144672	Sign	29048	2017-02-10 21:58	1
15007307	Order	-	2017-03-08 13:15	-
15007307	Consign	-	2017-03-10 17:14	-
15007307	Receive	49199	2017-03-14 19:27	-
15007307	Depart	49199	2017-03-14 19:51	-
15007307	Arrive	162115	2017-03-14 20:48	-
15007307	Depart	162115	2017-03-15 05:12	-
15007307	Arrive	166957	2017-03-15 06:29	-
15007307	Scan	166957	2017-03-15 07:28	-
15007307	Sign	166957	2017-03-15 10:04	5

Shipment 3144672 ended with an idle spell between the 6th and 10th of February, and shipment 15007307 began with an idle spell between the 10th and 14th of March. Cainiao’s operational transparency exposed this inactivity.

But the delivery service scores suggest that shipment 3144672’s late idleness was worse than shipment 15007307’s early idleness. This makes sense. First, imagine waiting for shipment 3144672: After seeing the package zip through three facilities in two days, you anticipate its arrival at any moment, only to suffer four additional days of delay. Moreover, the blustering start makes you more conscious of the subsequent silence—the parcel appears to vanish as its track-package signals abruptly end. Thus, when the delivery is finally made, you give it the worst possible score (1 out of 5).

Now imagine waiting for shipment 15007307: You see little progress in the first six days of your order. This is unsettling, but you’re not sure whether to attribute the lack of reported actions to a lack of reporting or to a lack of actions—only in the last two days do you learn that this shipper thoroughly records its activities. And by this time, you’re reassured by a steady stream of updates. This final hustle is fresh in your mind when you give the delivery the best possible score (5 out of 5).

These cherry-picked examples are outlandish, but they illustrate my thesis: consumers leave higher parcel delivery ratings when track-package activities occur near the final delivery time. Thus, the goodwill garnered by operational transparency depends on when the work is done—and when it’s not done.

I support this claim with Cainiao Network’s track-package records. Each shipment in my sample has a customer delivery score and a sequence of actions with corresponding timestamps. I regress

the delivery scores on the action times with five different specifications. Each indicates that later actions yield higher scores. For example, if the shipping time is 100 hours, then the first regression suggests that shifting a *single action* from hour 20 to hour 80 increases the expected delivery score by 0.021 standard deviations, and the second, third, fourth, and fifth regressions suggest that shifting the *average action* from hour 20 to hour 80 increases the expected delivery score by 0.075, 0.037, 0.185, and 0.197 standard deviations, respectively. For perspective, decreasing the shipping time from four to three days increases the expected delivery score by 0.064 standard deviations.

The first two regressions use OLS. The third regression uses pre-shipment delays to instrument for action times. Warehouse-to-shipper consignment is always the first action to follow the customer's order, so delaying this consignment delays all subsequent actions. But this consignment happens *before* the shipment, so the consignment time should not directly affect the shipment quality (conditional on the final delivery time). The fourth regression uses weekend lulls to instrument for action times. For example, since Saturdays and Sundays have the least activity, shipments that start on Fridays tend to have slower starts, and hence later average action times, whereas shipments that end on Mondays tend to have slower finishes, and hence earlier average action times. And the fifth regression generalizes this specification to exploit other temporal shocks, such as national holidays.

2 Theory

Delayed activity can increase scores in several ways. First, psychology's peak-end rule states that "the final moments of an extended episode appear to exert a strong influence on the overall judgment [of its utility]" (Varey and Kahneman, 1992, p. 169). For example, in Kahneman et al.'s (1993) peak-end study,

[s]ubjects were exposed to two aversive experiences: in the short trial, they immersed one hand in water at 14°C for 60 s; in the long trial, they immersed the other hand at 14°C for 60 s, then kept the hand in the water 30 s longer as the temperature of the water was gradually raised to 15°C, still painful but distinctly less so for most subjects. Subjects were later given a choice of which trial to repeat. A significant majority chose to repeat the long trial, apparently preferring more pain over less.

According to the peak-end rule, a shipment's ending will be especially memorable, which suggests that it's best to finish on a strong note with a burst of activity at the end.

Second, customer satisfaction depends on service quality *relative to expectations*. Surveying psychology's satisfaction literature, Oliver (1980, p. 460) found that

[a]lmost without exception, reviewers and early researchers in the areas of job, life, self, and patient satisfaction agree that satisfaction is a function of an initial standard and some perceived discrepancy from the initial reference point. ... Specifically, expectations are thought to create a frame of reference about which one makes a comparative judgment. Thus, outcomes poorer than expected (a negative disconfirmation) are rated

below this reference point, whereas those better than expected (a positive disconfirmation) are evaluated above this base.

In this light, early activity can be counterproductive, as it gives customers unrealistic expectations about the speed of delivery—starting fast raises customer hopes, and ending slow dashes them. Moreover, the unfulfilled expectations can make customers apprehensive, as Harvard’s Ryan Buell explains:

reading through your paper made me think of the work by Osuna (1985), which basically shows how customer uncertainty can increase frustration and anger, undermining people’s satisfaction. That’s the paper that basically became the reason we see progress bars everywhere—people value the certainty of knowing when a task will be complete or a service will be delivered. A package making fitful progress toward delivery ... only to be stalled at the last minute could amp up uncertainty.

Third, inactivity concerns customers only after they’ve learned to expect steady status updates. Most customers don’t know how much track-package activity to expect from a given shipper, so they can attribute a silent start to a silent shipper. But a lively start establishes a high benchmark: after a few days of consistent posting, a day of inactivity seems an ominous halt to momentum. Once trained to expect progress reports, customers will notice their absence.

And fourth, customers will check the track-package logs more frequently near the expected arrival time, so activities reported around this time are more likely to be noticed and appreciated.

3 Data

I use data provided for the 2018 MSOM Data Driven Research Challenge by Cainiao Network.¹ An affiliate of Alibaba Group, Cainiao runs an online logistics platform for managing the delivery of goods purchased through Alibaba’s websites. The company was founded in 2013 with the goal of “realiz[ing] delivery anywhere in China within 24 hours, and across the globe within 72 hours.”

The data set comprises (i) a 10.02 GB table that describes customer orders, (ii) a 507 MB table that describes warehouse inventories, (iii) a 2.52 GB table that describes products, (iv) a 77 MB table that describes merchants, and (v) a 74 GB table that describes package delivery logistics. The first table provides my dependent variable—the customer delivery score—and the last table provides my primary independent variables—the track-package action timestamps.

There are several types of track-package action:²

- Order: the customer places the order
- Consign: the warehouse dispatches the package

¹Cui et al. (2019), Li et al. (2019a), Li et al. (2019b), and Pendem and Deshpande (2019) also analyzed this data set as participants of the MSOM Data Driven Research Challenge. My work is closely related to Pendem and Deshpande’s (2019) paper, which studies the sales effect of the customer delivery scores, and is somewhat related to Cui et al.’s (2019) paper, which studies the sales effect of the temporary ban of a prominent shipper.

²There is also a confirmation action that I disregard, since it occurs after the customers file their delivery scores.



Figure 1: Example of What the Customer Sees

This is a screenshot from the Tmall mobile app that an anonymous reviewer graciously provided me. It lists a package's reported actions: the package was ordered on 12/12, consigned to the shipper on 12/13, moved from Handan to Xingtai on 12/13, moved from Xingtai to Shijiazhuang and then to Beijing on 12/14, and delivered on 12/15.

- Receive: the carrier receives the package
- Depart: the package exits a facility
- Arrive: the package enters a facility
- Scan: the shipper scans the package for final delivery
- Sign: the customer receives the package
- Failure: the shipper fails to deliver the package

The Tmall and GuoGuo mobile apps—which account for the lion's share of Cainiao's business—disclose these track-package actions in real time. For example, Figure 1 provides a screenshot of the actions Tmall reported to an anonymous reviewer as his or her package traveled from Handan to Beijing. However, this reviewer wanted me to tell you that he or she had to explicitly search out these records as neither Tmall nor GuoGuo notify customers when track-package actions are processed. Hence, only those curious enough to inspect the operation will observe the operation. (But a sizeable share should do so as it's quite easy to check the track-package records on the mobile apps.)

Finally, I filter my sample along several dimensions, removing all shipments

- with a failure action (0.74% of observations),
- with an origin warehouse not managed by Cainiao (73% of observations),
- without a shipment score or shipment times (64% of observations),

- with actions reported before the order action (0.024% of observations),
- with actions reported after the sign action (3.2% of observations),
- without exactly one sign action (2.6% of observations),
- without exactly one consign action (1.3% of observations),
- without the slowest shipping speed (15% of observations),
- with multiple shippers (0.0010% of observations),
- with multiple product types (6.7% of observations),
- with shipment times in excess of eight days (1.6% of observations),
- with more than ten posted actions (6.0% of observations), or
- with fewer than four posted actions (6.3% of observations).

The resulting sample comprises 101 thousand facilities, 4.68 million shipments, and 40.10 million actions from January 1, 2017 to July 31, 2017.

4 Variables

I now provide a glossary of the variables in my data. The most important variables are Delivery Score, Action Time, Average Action Time, and Action Count $[t, s]$.

- Delivery Score is a delivery logistics quality score left by the customer. The customer uploads this information via mobile app or website after receiving the package. The variable takes values in $\{1, \dots, 5\}$, where 1 is the worst and 5 the best, and has mean 4.82 and standard deviation 0.64 (see Table 2).
- Action Time is the time of a particular action, measured as a fraction of the shipping time. The variable takes values in $[0, 1]$, and has mean 0.49 and standard deviation 0.36 (see Table 3). For example, Table 1 reports shipment 3144672's Action Times: the order Action Time is 0.00 because the order action starts the shipment, the sign Action Time is 1.00 because the sign action ends the shipment, and the consign Action Time is 0.02 because the consign action happens after 2% of the shipping time has elapsed. Since order and sign Action Times mechanically equal 0.00 and 1.00, respectively, I henceforth disregard order and sign actions.
- Average Action Time is the shipment's average Action Time (excluding order and sign actions). This variable takes values in $[0, 1]$, and has mean 0.49 and standard deviation 0.13. For example, shipment 3144672's Average Action Time is $(0.020 + 0.028 + 0.054 + 0.104 + 0.112 + 0.136 + 0.139)/7 = 0.085$ (see Table 1).

- Action Count is the number of distinct actions—other than order and sign—reported on the shipment’s track-package log. This variable takes values in $\{4, \dots, 10\}$ (until §6 removes this constraint), and has mean 6.57 and standard deviation 1.72. For example, shipment 3144672’s Action Count is seven (see Table 1).
- Action Count $[t, s)$ is the number of distinct actions with Action Times in range $[t, s)$. This variable takes values in $\{0, \dots, 10\}$. For example, shipment 3144672 has an Action Count $[0, 0.05)$ of two (see Table 1).
- Shipping Time is the time between the shipment’s order and sign actions, measured in days. This variable takes values in $[0, 8]$ (until §6 removes this constraint), and has mean 2.74 and standard deviation 1.16. For example, shipment 3144672’s Shipping Time is 5.29 (see Table 1). A regression of Delivery Score on Shipping Time suggests that increasing the latter by one day decreases the former by 0.0443 points.
- Day Count is Shipping Time rounded up to the nearest day. This variable takes values in $\{1, \dots, 8\}$ (until §6 removes this constraint), and has mean 3.22 and standard deviation 1.19. For example, shipment 3144672’s Day Count is six (see Table 1).
- Consign Count, Receive Count, Arrive Count, Depart Count, and Scan Count are the number of distinct consign, receive, arrive, depart, and scan actions reported on the shipment’s track-package log. These variables take values in $\{0, \dots, 10\}$ (until §6 removes these constraints). For example, shipment 3144672 has a Receive Count of one and a Depart Count of two (see Table 1).
- Facility Count is the number of distinct facilities reported on the shipment’s track-package log.³ This variable takes values in $\{0, \dots, 8\}$, and has mean 3.52 and standard deviation 1.59. For example, shipment 3144672’s Facility Count is three (see Table 1).
- Day is the day of the shipment’s order action. This variable takes values in $\{1, \dots, 212\}$, where Day = 1 corresponds to 01/01/2017, the first date in my sample, and Day = 212 corresponds to 07/31/2017, the last date in my sample. For example, shipment 3144672’s Day is 36 since 02/05/2017 is the 36th date in my sample (see Table 1).
- Week is the week of the shipment’s order action. This variable takes values in $\{1, \dots, 31\}$, where Week = 1 corresponds to the week starting on 01/01/2017, the first Sunday in my sample, and Week = 31 corresponds to the week starting on 07/30/2017, the last Sunday in my sample. For example, shipment 3144672’s Week is 6, since 02/05/2017 falls in the sixth week of my sample (see Table 1).
- Day of Week is the day of the week of the shipment’s order action. This variable takes values in $\{1, \dots, 7\}$, where 1 corresponds to Sunday and 7 to Saturday. For example, shipment

³ Actually, Facility Count is the lesser of the number of reported facilities and eight. However, fewer than 0.02% of shipments have more than eight facilities.

Shipment	Action	Facility	Timestamp	Action Time
3144672	Order	–	2017-02-05 15:05	0.0000
3144672	Consign	–	2017-02-05 17:37	0.0200
3144672	Receive	105638	2017-02-05 18:40	0.0282
3144672	Depart	105638	2017-02-05 21:52	0.0535
3144672	Arrive	65132	2017-02-06 04:15	0.1038
3144672	Depart	65132	2017-02-06 05:20	0.1123
3144672	Arrive	29048	2017-02-06 08:22	0.1362
3144672	Scan	29048	2017-02-06 08:44	0.1391
3144672	Sign	29048	2017-02-10 21:58	1.0000

Table 1: Action Time Definition

I tabulate the Action Times of shipment 3144672. I derive these values from the action timestamps:

$$\text{Action Time}_n = \frac{\text{Timestamp}_n - \min_m(\text{Timestamp}_m)}{\max_m(\text{Timestamp}_m) - \min_m(\text{Timestamp}_m)}.$$

3144672's Day of Week is 1 since 02/05/2017 falls on a Sunday. Also, note that $\text{Day} = 7 \cdot (\text{Week} - 1) + \text{Day of Week}$.

- Buyer, Brand, Category, Merchant, and Shipper are the ID numbers of the customer, product brand, product category, merchant, and shipper. For example, shipment 3144672 has Buyer = 61581582.
- Shipping Speed is the shipping speed selected by the customer. This variable takes values in $\{1, 2, 3, \infty\}$, where the first three options guarantee delivery within 1, 2, and 3 days, respectively, and the last option provides no delivery date guarantee. Shipping Speeds are restricted to ∞ until §6.
- Post-Median Average Action Time is the average Action Time of the actions that occur after the median Action Time. This variable takes values in $[0, 1]$, and has mean 0.746 and standard deviation 0.138. For example, shipment 3144672's median Action Time is 0.1038, and its Post-Median Average Action Time is $(0.1123 + 0.1362 + 0.1391)/3 = 0.1292$ (see Table 1).

5 Results

Figure 2 demonstrates that shipments with different Delivery Scores have different Action Time distributions. The plots depict the Action Time probability density functions (PDFs) conditional on the Delivery Score minus the Action Time PDFs unconditional on the Delivery Score (I subtract away the unconditional distributions to highlight the across-score differences). Each action type yields the same pattern: to the left, the score-1–2 PDFs are the highest, followed by the score-3 PDFs, then the score-4 PDFs, and then the score-5 PDFs; to the right, this order is reversed. Thus, the score-1–2 actions occur earlier than the score-3 actions, which occur earlier than the score-4 actions, which occur earlier than the score-5 actions.

		Day Count								
		1	2	3	4	5	6	7	8	Total
Action Count	4	4.88	4.86	4.84	4.80	4.74	4.71	4.69	4.65	4.83
	5	4.89	4.86	4.83	4.80	4.75	4.71	4.68	4.68	4.82
	6	4.89	4.87	4.85	4.80	4.73	4.66	4.57	4.56	4.82
	7	4.88	4.86	4.83	4.80	4.75	4.69	4.64	4.59	4.82
	8	4.88	4.86	4.84	4.78	4.70	4.62	4.56	4.49	4.82
	9	4.88	4.85	4.83	4.78	4.72	4.65	4.62	4.53	4.81
	10	4.85	4.85	4.82	4.78	4.73	4.66	4.61	4.53	4.79
Total		4.88	4.86	4.84	4.79	4.74	4.68	4.63	4.59	4.82

Table 2: Average Delivery Scores

I tabulate the average Delivery Score, by Action Count and Day Count. For example, four-action shipments that arrive within a day have an average Delivery Score of 4.88.

	10%	20%	30%	40%	50%	60%	70%	80%	90%
Order	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Consign	0.03	0.04	0.06	0.09	0.13	0.18	0.22	0.27	0.35
Receive	0.07	0.10	0.14	0.19	0.23	0.28	0.33	0.40	0.51
Depart	0.17	0.27	0.35	0.43	0.52	0.59	0.67	0.74	0.83
Arrive	0.23	0.36	0.48	0.57	0.65	0.73	0.81	0.86	0.92
Scan	0.69	0.81	0.85	0.88	0.91	0.93	0.95	0.97	0.99
Sign	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00

Table 3: Action Time Deciles

I tabulate the Action Time deciles, by action type. For example, the median consign action happens after 13% of the shipping time has elapsed. Note that the order Action Time is always 0.00 and the sign Action Time is always 1.00 because order and sign actions bookend the shipment.




Figure 2: Action Time Distributions

I estimate each action type's Action Time PDFs, both conditional and unconditional on the Delivery Score. I then subtract the unconditional density estimates from the conditional density estimates, and I plot the differences' 90% confidence intervals with lines of varying thickness, thinner lines indicating more precise estimates. For example, the score-5 lines are the thinnest because score-5 estimates are the most precise, and the score-5 estimates are the most precise because 89.9% of scores are 5. Since so many scores are 5, the distributions conditional on the score being 5 resemble the unconditional distributions, which explains why the score-5 lines are so near zero. I combine the score-1 and score-2 observations because only 0.5% of scores are 2.

Merchant	170
Brand	98
Category	15
Shipper	20
Week	30
Day Count	7
Facility Count	8
Arrive Count	6
Depart Count	7
Receive Count	5
Scan Count	8

Table 4: Control Variable Dummies

This table records the number of dummy variables of each type that I use as controls in my regressions. For Day Count, Facility Count, Receive Count, Arrive Count, Depart Count, and Scan Count, the number of dummy variables equals one less than the number of distinct values (the fully saturated case). For Brand, Category, Merchant, Shipper, and Week, there is one dummy variable for each value with at least 5,000 observations. For example, the sample has 186 Brand values, 98 of which appear at least 5,000 times, and eight Day Count values.

I now formalize the relationship between Action Times and Delivery Scores with a set of regressions. I run 14 OLS regressions across 14 subsamples, slicing the data by Action Count and Day Count. The dependent variable is the Delivery Score. The control variables are a set of dummies that specify (i) the values of Brand, Category, Merchant, Shipper, and Week that have at least 5,000 observations, and (ii) all the values of Day Count, Facility Count, Receive Count, Arrive Count, Depart Count, and Scan Count (see Table 4).⁴ Finally, the primary independent variables are the Action Count[t, s] values corresponding to the following 19 time ranges:

$$\begin{aligned}
& [0.00, 0.05), [0.05, 0.10), [0.10, 0.15), [0.15, 0.20), [0.20, 0.25), \\
& [0.25, 0.30), [0.30, 0.35), [0.35, 0.40), [0.40, 0.45), [0.45, 0.50), \\
& [0.50, 0.55), [0.55, 0.60), [0.60, 0.65), [0.65, 0.70), [0.70, 0.75), \\
& [0.75, 0.80), [0.80, 0.85), [0.85, 0.90), \quad \text{and} \quad [0.90, 0.95).
\end{aligned}$$

For example, the Action Times of shipment 3144672 are

$$\begin{array}{ccccccc}
0.020, & 0.028, & 0.054, & 0.104, & 0.112, & 0.136, & \text{and } 0.139, \\
\underbrace{\hspace{1.5cm}}_{\in [0.00, 0.05)} & \underbrace{\hspace{1.5cm}}_{\in [0.05, 0.10)} & \underbrace{\hspace{1.5cm}}_{\in [0.10, 0.15)} & & & &
\end{array}$$

so for this observation, Action Count[0.00, 0.05) is two, Action Count[0.05, 0.10) is one, Action Count[0.10, 0.15) is four, and the other primary independent variables are zero.

Figure 3 plots the 19 time ranges' 19 coefficient estimates.⁵ These estimates report the amount an action in the given time range increases the expected score minus the amount an action in the [0.95, 1.00] time range increases the expected score. For example, the far-left estimate of the

⁴My controls do not include Consign Count dummies because Consign Count always equals one (for now).

⁵I calculate all standard errors in this work with the paired bootstrap, which is robust to general heteroskedasticity (Cameron and Trivedi, 2005, p. 376).

four-action plot is -0.017, which suggests that shifting one Action Time from [0.00, 0.05) to [0.95, 1.00] increases the expected Delivery Score by 0.017 points (or 0.027 standard deviations). Overall, the estimates suggest that actions that occur after 95% of the shipping time has elapsed increase scores more than

- actions that occur before 10% of the shipping time has elapsed: of the 28 Action Count[0.00, 0.05)–Action Count[0.05, 0.10) estimates, 27 are negative and 21 are significantly negative at the $p = 0.01$ level;
- actions that occur before 25% of the shipping time has elapsed: of the 70 Action Count[0.00, 0.05)–Action Count[0.20, 0.25) estimates, 67 are negative and 51 are significantly negative at the $p = 0.01$ level; and
- actions that occur before 50% of the shipping time has elapsed: of the 140 Action Count[0.00, 0.05)–Action Count[0.45, 0.50) estimates, 136 are negative and 99 are significantly negative at the $p = 0.01$ level.

However, the most valuable actions appear to be those that occur between Action Times 0.85 and 0.95: of the 28 Action Count[0.85, 0.90)–Action Count[0.90, 0.95) estimates, 24 are positive and 7 are significantly positive at the $p = 0.01$ level. The average difference between the Action Count[0.80, 0.85) and Action Count[0.15, 0.20) estimates is 0.0137, which suggests that shifting one action from the [0.15, 0.20) time range to the [0.80, 0.85) time range increases the expected Delivery Score by an average of 0.0137 points (which is quite a lot, considering that the average shipment comprises 6.57 actions).

The [0.85, 0.95) time range is the sweet spot because it's late enough to enjoy a peak-end effect but not so late that the package arrives before the action is noticed. Note, customers aren't notified when actions are uploaded, so many actions posted in the [0.85, 0.95) time range won't be noticed until the [0.95, 1.00) time range, and many actions posted in the [0.95, 1.00) time range won't be noticed until they're moot. However, despite this time lag, the correlation between Action Times and Delivery Scores is undeniably positive: 13 out of 14 trend lines fit through the regression estimates are significantly positive at the $p = 0.01$ level.

To establish the ubiquity of this effect, I run 100 additional regressions across 50 different subsamples. The subsamples are the observations of the ten most frequent (i) Brand, (ii) Category, (iii) Merchant, (iv) Shipper, and (v) Week values. For example, the 10 most common Shipper IDs are 247, 565, 674, 724, 532, 431, 149, 132, 270, and 184, each of which has its own subsample. For each subsample, I run two regressions. For both regressions, the dependent variable is the Delivery Score and the primary independent variable is the Average Action Time, but the first regression incorporates Table 4's control variables, while the second does not.

Table 5 reports the Average Action Time coefficient estimates. Of the 100 estimates, 93 are positive and 84 are significantly positive at the $p = 0.01$ level: the effect is pervasive.⁶ And the

⁶ Suppose we adopt the null hypothesis that the Average Action Time has no effect on the Delivery Score. In this case, the number of estimates, out of 100, that are significantly positive at the $p = .01$ level would have a binomial(100, .01) distribution (ignoring dependencies across regressions). Under this null hypothesis, the probability of at least

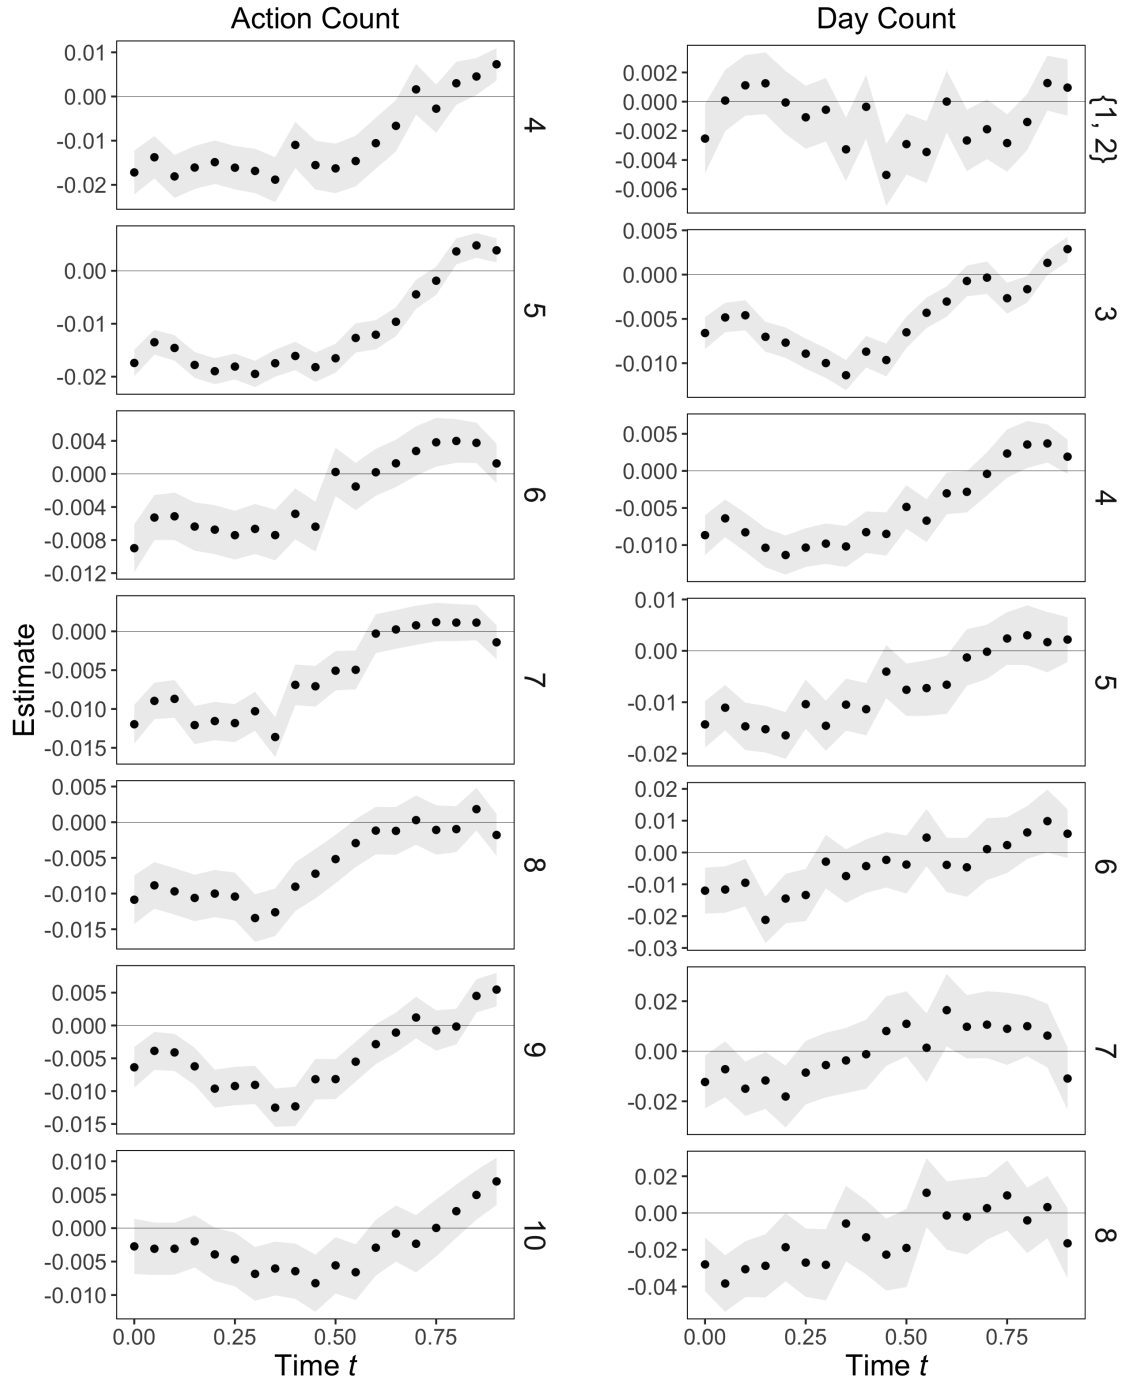


Figure 3: OLS Estimates of Delivery Score on Action Count $[t, s]$

I run 14 regressions across 14 subsamples: the observations with Action Count = n , for $n \in \{4, \dots, 10\}$, and the observations with Day Count $\in n$, for $n \in \{\{1, 2\}, 3, \dots, 8\}$. I pool the Day Count = 1 and Day Count = 2 observations because only 2% of Day Counts are 1. The dependent variable is Delivery Score; the control variables are Table 4's dummies; and the primary independent variables are Action Count $[t, t + 0.05)$, for $t \in \{0.00, 0.05, \dots, 0.90\}$. The black dots depict the coefficient estimates of these primary independent variables, with the left-most dot corresponding to the $[0.00, 0.05)$ time range and the right-most dot corresponding to the $[0.90, 0.95)$ time range. The gray bands depict the estimates' 90% confidence intervals.

effect is meaningful: Running the regressions across the entire sample yields an estimate of 0.081 (see Table 6), which suggests that shifting the Average Action Time from 0.2 to 0.8 increases the expected Delivery Score by $0.081 \cdot (.8 - .2) = 0.049$ points (or 0.075 standard deviations). This change would have the same effect on Delivery Scores as a $0.049/0.0443 = 1.11$ -day reduction in Shipping Times (recall that shortening the Shipping Time by one day increases the expected Delivery Score by 0.0443 points).

To establish the causality of the relationship between Average Action Time and Delivery Score, I run three sets of two-stage least squares (2SLS) regressions with three sets of instrumental variables (IV). These IV regressions are analogous to those reported in Table 5, except they permit the Average Action Time to correlate with the error term via unobserved shipping factors. For example, suppose the final deliveryman is either tardy and rude or prompt and courteous; in this case, packages with late delays would tend to arrive in a ruder fashion than those with early delays. To control for such unobserved shipping factors, I use instrumental variables that shift the Average Action Time yet remain independent of the shipping process.

My first 2SLS specification derives instruments from the weekly variation in activity levels. For example, there's more idleness on weekends, so the idleness of Saturday-to-Friday shipments is earlier than the idleness of Monday-to-Sunday shipments; thus, Saturday-to-Friday shipments have larger Average Action Times than Monday-to-Sunday shipments. This logic suggests that I can treat Day of Week and Day Count pairs as exogenous shifters of Average Action Time. I interact these pairings with Table 4's Shipper dummies, since different companies have different weekly trends. My resulting instrumental variables are 1,159 Day of Week \times Day Count \times Shipper dummies (plus the control variables when the specification has them). For example, shipment 3144672 starts on the first day of the week, ends on the sixth day, and is handled by Shipper 149, so for this observation the $\{\text{Day of Week} = 1\} \times \{\text{Day Count} = 6\} \times \{\text{Shipper} = 149\}$ dummy variable is one and the other 1,158 dummy variables are zero. These instruments explain 20.9% of the variation in Average Action Times.

My second 2SLS specification is the same as the first except it replaces Day of Week with Day. That is, it uses 25,617 Day \times Day Count \times Shipper dummies as instruments (plus the control variables when the specification has them). Ryan Buell, at Harvard, and an anonymous reviewer gave me the idea for these instrumental variables; they explained that Day of Week isn't granular enough to capture most temporal shocks, such as national holidays, inclement weather, or site-wide promotions. Giving each day its own set of instruments enables me to more flexibly exploit calendar events. These instruments explain 32.5% of the variation in Average Action Times.

My final 2SLS specification instruments for the Average Action Time with the consign Action Time. From the customer's perspective, it appears that the shipment starts as soon as the order is placed. But, actually, the shipment doesn't begin until the warehouse consigns the parcel to the shipper. Since the consign Action Time should not directly affect the condition of the package, the consign Action Time should not directly affect the Delivery Score (after controlling for the package's arrival time). I therefore treat the warehouse-to-shipper consignment times as exogenous

12 significantly positive estimates is only $4.7 \cdot 10^{-10}$. Thus, with our 84 significantly positive estimates, we strongly reject the null.

		$n = 1$	$n = 2$	$n = 3$	$n = 4$	$n = 5$	$n = 6$	$n = 7$	$n = 8$	$n = 9$	$n = 10$
Controls	Brand	0.067 (0.006)	0.039 (0.010)	0.119 (0.013)	0.013 (0.013)	0.137 (0.015)	-0.036 (0.014)	0.114 (0.011)	0.069 (0.015)	0.093 (0.018)	0.084 (0.017)
	Category	0.078 (0.006)	0.041 (0.005)	0.036 (0.008)	0.119 (0.006)	0.043 (0.008)	0.063 (0.013)	0.086 (0.017)	0.068 (0.016)	0.074 (0.020)	0.017 (0.025)
	Merchant	0.063 (0.008)	0.119 (0.013)	0.010 (0.015)	0.137 (0.017)	0.031 (0.020)	-0.035 (0.016)	0.032 (0.021)	0.077 (0.021)	-0.059 (0.022)	0.042 (0.019)
	Shipper	0.079 (0.007)	0.049 (0.008)	0.013 (0.008)	0.119 (0.006)	0.109 (0.011)	0.048 (0.011)	0.179 (0.010)	0.111 (0.017)	0.052 (0.019)	-0.052 (0.018)
	Week	0.120 (0.008)	0.071 (0.010)	0.064 (0.014)	0.050 (0.015)	0.122 (0.015)	0.046 (0.015)	0.070 (0.014)	0.068 (0.017)	0.047 (0.015)	0.046 (0.017)
No Controls	Brand	0.114 (0.005)	0.111 (0.008)	0.298 (0.011)	0.230 (0.010)	0.167 (0.012)	-0.021 (0.012)	0.099 (0.010)	0.099 (0.012)	0.114 (0.015)	0.169 (0.013)
	Category	0.174 (0.005)	0.122 (0.005)	0.166 (0.007)	0.110 (0.006)	0.115 (0.007)	0.177 (0.010)	0.219 (0.013)	0.089 (0.013)	0.096 (0.017)	0.158 (0.021)
	Merchant	0.097 (0.008)	0.298 (0.011)	0.226 (0.012)	0.161 (0.014)	0.130 (0.016)	-0.045 (0.015)	0.114 (0.017)	0.113 (0.017)	-0.019 (0.019)	0.129 (0.016)
	Shipper	0.187 (0.006)	0.189 (0.006)	0.005 (0.007)	0.108 (0.006)	0.274 (0.010)	0.034 (0.009)	0.290 (0.010)	0.425 (0.014)	0.315 (0.017)	0.028 (0.016)
	Week	0.100 (0.007)	0.127 (0.009)	0.163 (0.013)	0.190 (0.012)	0.263 (0.013)	0.122 (0.012)	0.118 (0.012)	0.183 (0.014)	0.115 (0.013)	0.108 (0.014)

Table 5: OLS Estimates of Delivery Score on Average Action Time

I run two regression specifications across 50 subsamples, for a total of 100 regressions. The subsamples are the observations with the n th most common Brand, Category, Merchant, Shipper, and Week values, for $n \in \{1, \dots, 10\}$. The first regression specification includes Table 4's control variables, and the second does not. The dependent variable is Delivery Score, and the primary independent variable is Average Action Time. I tabulate the Average Action Time coefficient estimates and corresponding standard errors.

shifters of Average Action Time. Specifically, I create a set of dummy variables that characterize the consign Action Time's decile. For example, the consign Action Time is in the first decile if less than 0.024, is in the second decile if between 0.024 and 0.039, and is in the third decile if between 0.039 and 0.057. I then interact these consign Action Time decile dummies with Table 4's Shipper dummies, since the relationship between consign Action Time and Average Action Time should vary by shipper. My resulting instrumental variables are 209 consign Action Time decile \times Shipper dummies (plus the control variables when the specification has them). For example, shipment 3144672 has a first-decile consign Action Time and a Shipper ID of 149, so for this observation the $\{\text{consign Action Time decile} = 1\} \times \{\text{Shipper} = 149\}$ dummy variable is one and the other 208 dummy variables are zero. These instruments explain 62.2% of the variation in Average Action Times.

Table 6 reports the 2SLS estimates. They are all significantly positive at the $p = 0.001$ level, and are similar to the corresponding OLS estimates.

	Controls	No Controls
OLS	0.081 (0.003)	0.137 (0.002)
2SLS: Day of Week	0.200 (0.016)	0.317 (0.005)
2SLS: Day	0.213 (0.007)	0.289 (0.004)
2SLS: Consign Action Time	0.040 (0.003)	0.119 (0.003)

Table 6: OLS and 2SLS Estimates of Delivery Score on Average Action Time
I run two OLS regressions and six 2SLS regressions. The OLS regressions are the same as Table 5's, except they use the entire sample. The 2SLS regressions are the same as the OLS regressions, except they instrument for the Average Action Time. The first 2SLS specification uses 1,159 Day of Week \times Day Count \times Shipper dummies as instruments, the second uses 25,617 Day \times Day Count \times Shipper dummies as instruments, and the third uses 209 consign Action Time decile \times Shipper dummies as instruments. Additionally, the control variables serve as exogenous instruments (in the regressions that include them). I tabulate the Average Action Time coefficient estimates and corresponding standard errors.

6 Replication

I imposed 14 filters on my data in §3. I did so to compare like with like and to minimize the effect of unobserved confounding variables: e.g., if a shipment requires three weeks to arrive, then there's probably something important about that delivery I don't see. But Bill Schmidt at Cornell objected to some of these restrictions.⁷ To address his feedback, I reran my regressions with the observations initially left out. Loosely speaking, this analysis provides a replication of my primary results because I conducted it with a different sample after distributing §5's findings.

To create my replication sample, I begin with the observations I excluded in §3 and remove shipments

- with a failure action (0.74% of observations),
- with an origin warehouse not managed by Cainiao (73% of observations),
- without a shipment score or shipment times (64% of observations),
- with actions reported before the order action (0.024% of observations),
- with actions reported after the sign action (3.2% of observations), or
- without exactly one sign action (2.6% of observations).

⁷By the way, Bill Schmidt has recently written on operational transparency with Ananth Raman. Schmidt and Raman (2018) show that operational transparency can decrease the information asymmetry between a company and its investors, which in turn can make the company's stock price less sensitive to operational disruptions.

	1 Day	2 Days	3 Days	4+ Days
Delivery Score	4.90	4.88	4.85	4.78
Action Time	0.50	0.48	0.51	0.48
Average Action Time	0.50	0.47	0.51	0.45
Shipping Time	0.53	1.20	1.92	3.72
Action Count	4.68	5.51	5.96	7.23
Facility Count	2.71	2.81	3.03	4.19

Table 7: Replication Sample Summary Statistics

I tabulate the average of six variables in my replication sample by Shipping Speed. The 1-day Shipping Speed is the fastest, with an overnight delivery guarantee, and the ∞ -day Shipping Speed is the slowest, with no guaranteed delivery date.

In other words, my replication sample comprises the observations that satisfy the first six conditions of §3 but not the last seven. It comprises 7.78 million shipments and 66.4 million actions, none of which appear in my initial sample, and it contains a new variable: Shipping Speed, which I initially restricted to the slowest setting. Table 7 demonstrates that slowing the Shipping Speed from 1 day to 2 days, to 3 days, to ∞ days increases the Shipping Time and Action Count and decreases the Delivery Score.⁸

I rerun Figure 3’s regression with my replication sample and plot the estimates in Figure 4. The effect is stronger when the Shipping Speed is slower: from 1 day to 2 days to 3 days to ∞ days, the slope of the trend line through the estimates increases from 0.0025 to 0.0104 to 0.0137 to 0.0209. This makes sense: strengthening the shipping guarantee shortens the Shipping Time, which blurs the distinction between early and late actions (e.g., Figure 3 finds no significant results for Shipping Times shorter than two days).

Next, I rerun Table 6’s regressions with my replication sample and tabulate the coefficient estimates in Table 8. As before, the effect is stronger with slower Shipping Speeds: from 1 day to 2 days to 3 days to ∞ days, the primary OLS estimate increases from -0.001 to 0.043 to 0.068 to 0.143. All the 2-, 3-, and ∞ -day estimates are significantly positive, and three-quarters of the 1-day estimates are significantly positive.

7 Mechanism

In §2, I explained that delaying actions could increase scores by (i) making actions more memorable; (ii) giving customers more conservative arrival time estimates; (iii) making customers believe that actions go unreported; and (iv) making actions occur at more conspicuous times. Of the four, I believe the first driver—the peak-end effect—is most influential.⁹

⁸In addition to multiple Shipping Speeds, there are now shipments with multiple shippers, product types, and consign actions. To accommodate these changes, I (i) run my regressions by Shipping Speed, (ii) define Shipper as the first shipper to handle the package, (iii) define Brand and Category as the brand and category of the first listed product type, (iv) include Consign Count dummies as control variables, and (v) derive the consign Action Time decile \times Shipper instrumental variables from the time of the first consign action.

⁹Many thanks to Ryan Buell, at Harvard, who gave me most of the ideas for this section.

		1 Day	2 Days	3 Days	∞ Days
Controls	OLS	-0.001 (0.006)	0.043 (0.002)	0.068 (0.003)	0.143 (0.007)
	2SLS: Day of Week	0.288 (0.064)	0.141 (0.017)	0.195 (0.012)	0.376 (0.035)
	2SLS: Day	0.090 (0.016)	0.154 (0.005)	0.170 (0.007)	0.193 (0.009)
	2SLS: Consign Action Time	-0.017 (0.006)	0.031 (0.003)	0.036 (0.004)	0.174 (0.005)
No Controls	OLS	0.057 (0.008)	0.094 (0.002)	0.145 (0.003)	0.137 (0.010)
	2SLS: Day of Week	0.718 (0.023)	0.414 (0.005)	0.371 (0.004)	0.214 (0.033)
	2SLS: Day	0.372 (0.017)	0.263 (0.005)	0.288 (0.005)	0.162 (0.006)
	2SLS: Consign Action Time	0.056 (0.008)	0.091 (0.004)	0.128 (0.003)	0.229 (0.005)

Table 8: Replication Estimates of Delivery Score on Average Action Time
I implement Table 6's regressions with my replication sample. To accommodate the new sample, I add a Consign Count control variable and slightly modify the consign Action Time decile \times Shipper instrumental variables and the Shipper, Brand, and Category control variables. I also run the regressions by Shipping Speed.




Figure 4: Replication Estimates of Delivery Score on Action Count $[t, s]$

I implement Figure 3's regression with my replication sample. To accommodate the new sample, I add a Consign Count control variable and slightly modify the Shipper, Brand, and Category control variables. I also run the regression by Shipping Speed rather than by Action Count and Day Count.

First, only 1.46% of shipments with 2- or 3-day Shipping Speeds arrive late, so customers should have firm, predetermined, and accurate beliefs about when these shipments will arrive. Thus, it's difficult to attribute Figure 4's results to the second driver, expectations management.

Second, Figure 3's curves increase convexly. Fitting each set of estimates to a second-order polynomial, I find seven out of the 14 quadratic terms significantly positive at the $p = .01$ level, and zero significantly negative. To establish this convexity more formally, I rerun Table 6's OLS regressions with Post-Median Average Action Time as an additional regressor. The Post-Median coefficient estimates are significantly positive (see Table 9), which suggests that Delivery Scores are more sensitive to actions that occur after the median action. This fact—that later actions are more influential—is consistent with the first driver, but inconsistent with the second and third drivers. Indeed, if our results stemmed from customers setting overly optimistic expectations or attributing silence to non-reporting, then the earlier actions would be more influential, because it is these actions that have the most influence over the customers' arrival-time and status-update-frequency beliefs. Accordingly, if the second or third drivers were responsible, then we would expect Figure 3's curves to be concave and Table 9's Post-Median Average Action Time coefficients to be negative.

Finally, Table 10 shows that the results hold in the subset of shipments that arrived more than 24 hours early. To create this table, I reran Table 8's OLS regressions across (i) the subsample of observations with 2-day Shipping Speeds and less-than-one-day Shipping Times and (ii) the subsample of observations with 3-day Shipping Speeds and less-than-two-day Shipping Times. Since they arrived so far ahead of schedule, these shipments should not have arrived at a particularly

	Controls	No Controls
Average Action Time	0.036 (0.004)	0.034 (0.004)
Post-Median Average Action Time	0.054 (0.004)	0.120 (0.004)

Table 9: OLS Estimates of Delivery Score on Post-Median Average Action Time
I implement Table 6's OLS regressions with an additional covariate: the Post-Median Average Action Time. The Post-Median estimates report the difference between the pre-median Action Time effect on Delivery Scores and the post-median Action Time effect on Delivery Scores. For example, the leftmost column suggests that increasing a pre-median Action Time by 0.1 would increase the expected Delivery Score by $0.1 \cdot 0.036 = 0.0036$ points, whereas increasing a post-median Action Time by 0.1 would increase the expected Delivery Score by $0.1 \cdot (0.036 + 0.054) = 0.0090$ points.

	Controls	No Controls
2 Days	0.0060 (0.0034)	0.0172 (0.0029)
3 Days	0.0269 (0.0035)	0.0228 (0.0031)

Table 10: OLS Estimates from Shipments that Arrive Ahead of Schedule
I apply Table 8's OLS regressions to the subset of shipments that arrived at least one day early. The top row corresponds to observations with Shipping Speed = 2 and Day Count = 1, and the bottom row corresponds to observations with Shipping Speed = 3 and Day Count ≤ 2 .

conspicuous time. Thus, it's difficult to attribute Table 10's results to the fourth driver.

8 Robustness Checks

I now run three robustness checks with my initial sample.

First, Dennis Zhang at Washington University and an anonymous reviewer identified a potential problem: both action rates and service quality vary by location. For example, my estimates would be biased upwards if packages moved more expeditiously through the city than through the country and urban customers left systematically higher scores than rural customers. To control for geographic effects, I match my sample by the final facility reported on the track-package log (with 104,000 distinct locations, this final facility variable is granular). I consider two matching specifications: the first randomly pairs shipments with the same final facility, and the second randomly pairs shipments with the same final facility, Brand, Category, and Merchant. After pairing the observations, I difference each pair's Delivery Scores and Average Action Times, and regress the differenced Delivery Scores on the differenced Average Action Times. Table 11 demonstrates that controlling for the final facility location does not overturn the result.

Second, Ruomeng Cui at Emory University identified a second potential problem: customers can cancel a shipment at no cost before the consign action. This can introduce a selection bias, as the time until the first action influences whether the transaction is represented in my sample (which does not include canceled shipments).¹⁰ To avoid this potential bias, I control for the time until the consign action (measured in hours, not as a fraction of the Shipping Time). I group the consign times into 100 percentile buckets and match the sample by bucket. I consider two matching specifications: the first randomly pairs shipments with the same consign time bucket, and the second randomly pairs shipments with the same consign time bucket, Brand, Category, and Merchant. As before, I difference the data across pairs and regress the differenced Delivery Scores on the differenced Average Action Times. Table 11 demonstrates that controlling for the consignment time does not overturn the result.

Third, an anonymous reviewer asked me to control for the identity of the customer. I control for customer identity with two matching specifications: the first randomly pairs shipments with the same Buyer, and the second randomly pairs shipments with the same Buyer, Brand, Category, and Merchant. As before, I difference the data across pairs and regress the differenced Delivery Scores on the differenced Average Action Times. Table 11 demonstrates that controlling for the customer identity does not overturn the result.

Fourth, an anonymous reviewer wondered whether I could reproduce my result without the Delivery Score ≤ 2 observations. Thus, I rerun Table 6's OLS regressions without these extreme observations. Without control variables, I get an Average Action Time coefficient estimate of 0.021, with a corresponding t-statistic of 12.2; with control variables, I get a coefficient estimate of 0.051, with a corresponding t-statistic of 40.5. So I don't need extreme Delivery Scores to get a significant result.

9 Conclusion

I show that Cainiao's customers leave higher delivery scores when the track-package activities they see gravitate toward the end of the shipping horizon. Figure 3's estimates suggest that increasing *one* action time from [0.10, 0.15) to [0.80, 0.85) increases the expected delivery score by an average of 0.0141 points, and Table 6's primary OLS estimates suggest that increasing the *average* action time from 0.15 to 0.85 increases the expected delivery score by 0.0565 points. On average, an extra day of shipping decreases the expected delivery score by 0.0443 points, so these interventions are roughly analogous to decreasing the shipping time by $0.0141/0.0443 = 0.318$ and $0.0565/0.0443 = 1.28$ days.

These results are consistent with the peak-end rule, which states that customers remember endings more vividly than beginnings. Accordingly, Cainiao should emphasize last-mile logistics, as the last mile is the most memorable mile. Alternatively, Cainiao could craft their messages to

¹⁰Incidentally, Dennis Zhang and Ruomeng Cui have recently written an operational transparency paper with Achal Bassamboo. Cui et al. (2018) exogenously shifted the inventory levels posted in Amazon Lightning Deals by randomly adding products to 10 fictitious Amazon carts. Cui et al. (2018, p. 16) showed that reducing the available inventory levels increased demand rates, concluding that "real-time inventory information could serve as an effective lever for signaling popularity and attracting future customers."

	Controls	No Controls
Final Facility	0.103 (0.004)	0.092 (0.004)
Consign Time	0.220 (0.004)	0.190 (0.004)
Buyer	0.037 (0.011)	0.047 (0.006)

Table 11: Robustness Check Estimates

I run six regressions, each with its own matching scheme. For a given scheme, I create random matched pairs and difference the Delivery Score and Average Action Time variables by pair. I then regress the differenced Delivery Score variable on an intercept and the differenced Average Action Time variable, reporting the coefficient estimates of the latter. The “no controls” specifications match the shipments by either the final facility reported on the track-package logs, the consign time percentile, or the Buyer. The “controls” specifications match the shipments by these variables in addition to Brand, Category, and Merchant.

highlight later actions—e.g., not reporting the consign actions would increase the average action time from 0.494 to 0.553.

The peak-end rule should apply to most service operations. For example, Redelmeier et al. (2003) showed that adding a needless resting period to the end of a colonoscopy improved patient impressions of the procedure:

By random assignment, half the patients had a short interval added to the end of their procedure during which the tip of the colonoscope remained in the rectum. ... As theorized, patients who underwent the extended procedure experienced the final moments as less painful (1.7 vs. 2.5 on a ten point intensity scale, $P < 0.001$), rated the entire experience as less unpleasant (4.4 vs. 4.9 on a 10 cm visual analogue scale, $P = 0.006$), and ranked the procedure as less aversive compared to seven other unpleasant experiences (4.1 vs. 4.6 with eight as the worst, $P = 0.002$).

This experiment was published in a medical journal, but at its core the work is operations management: the researchers modify a repeated process (medical procedure) to reduce its perceived cost (recalled pain). The intervention challenges our operational insights—lengthening the procedure increases the flow time and decreases the throughput rate—and our basic intuitions—lengthening the procedure increases the total experienced pain. But in this case, the relevant quality measure is not the experienced pain, but the recollected pain, as “Patients’ memories of unpleasant medical procedures influence their decisions about future treatment choices” (Redelmeier et al., 2003, p. 187). And a patient’s retrospective evaluation is not a naïve integral of instantaneous utilities, since the final utilities receive extra weight. Accordingly, the researchers modify the process to end on a (relatively) painless note. This all’s-well-that-ends-well logic applies more broadly: judgment usually comes after the process, not during the process. This means we should pay special attention to how our service operations end.

That said, here's my attempt at an agreeable ending: What's more frustrating, waiting six weeks for the referees and one week for the editor, or waiting one week for the referees and six weeks for the editor?

References

- Buell, Ryan W., Tami Kim, Chia-Jung Tsay. 2017. Creating Reciprocal Value Through Operational Transparency. *Management Science* **63**(6) 1673–1695.
- Buell, Ryan W., Michael I. Norton. 2011. The labor illusion: How operational transparency increases perceived value. *Management Science* **57**(9) 1564–1579.
- Cameron, A. C., P. K. Trivedi. 2005. *Microeconometrics: Methods and Applications*. Cambridge University Press, Cambridge, England.
- Cui, Ruomeng, Meng Li, Qiang Li. 2019. Value of High-Quality Logistics: Evidence From a Clash Between SF Express and Alibaba. *Working Paper* (January 2018) 1–37.
- Cui, Ruomeng, Dennis J Zhang, Achal Bassamboo. 2018. Learning from Inventory Availability Information: Evidence from Field Experiments on Amazon. *Management Science* (Articles in Advance).
- Kahneman, Daniel, Barbara L. Fredrickson, C. A. Schreiber, D. A. Redelmeier. 1993. When more pain is preferred to less. *Psychological Science* **4** 401–405.
- Li, Maggie, Xiang Liu, Yan Huang, Cong Shi. 2019a. Integrating Empirical Estimation and Assortment Personalization for E-Commerce: A Consider-then-Choose Model. *Working Paper* 1–38.
- Li, Xiaocheng, Yufeng Zheng, Zhenpeng Zhou, Zeyu Zheng. 2019b. Demand Prediction, Predictive Shipping, and Product Allocation for Large-scale E-commerce. *Working Paper* 1–32.
- Oliver, Richard L. 1980. A Cognitive Model of the Antecedents and Consequences of Satisfaction Decisions. *Journal of Marketing Research* **17**(4) 460–469.
- Osuna, Edgar Elías. 1985. The psychological cost of waiting. *Journal of Mathematical Psychology* **29**(1) 82–105.
- Pendem, Pradeep, Vinayak Deshpande. 2019. Logistics Performance, Ratings, and its Impact on Customer Purchasing Behavior and Sales in E-commerce Platforms. *Working Paper* 1–32.
- Redelmeier, Donald A., Joel Katz, Daniel Kahneman. 2003. Memories of colonoscopy: A randomized trial. *Pain* **104**(1-2) 187–194.
- Schmidt, William, Ananth Raman. 2018. Operational Disruptions, Firm Risk, and Control Systems. *Working Paper* .
- Varey, Carol, Daniel Kahneman. 1992. Experiences extended across time: Evaluation of moments and episodes. *Journal of Behavioral Decision Making* **5**(3) 169–185.